

REVIEW

★★★★★ Project story with exact ownership

★★★★★ One production-failure or scale scenario

★★★★★ RAG or fine-tuning trade-offs

PRACTICE OUT LOUD

★★★★★ 90-sec project walkthrough

★★★★★ A coding problem out loud, narrating complexity

★★★★★ One Googliness-style behavioral story

REMEMBER

✓ State your exact contribution, not the team's

✓ Clarify assumptions before coding

✓ Quantify impact wherever possible

ROUND COVERAGE / TOP TOPICS

candidate-level, out of 10

TOP QUESTIONS

rehearse these

★★★★■	Coding / DSA (algorithm problem, complexity analysis)	7/10
★★★★■	Production ML reasoning (deploy, debug, scale a live system)	7/10
★★■	Current AI/ML project deep dive, own exact contribution	5/10
★★■	LLM / GenAI reasoning (RAG, fine-tuning, inference, agents)	6/10
★★■	RAG / retrieval / vector search	4/10
★★■	End-to-end system design (not just a coding problem)	4/10
★★■	Model evaluation / monitoring / drift	4/10
★★■	Behavioral / "Googliness" leadership round	2/10

OBSERVED RESUME FOCUS

- ★★★★★ Current AI/ML project and exact personal ownership
- ★★★★★ Production deployment + measured impact, not just model accuracy
- ★★★★■ Architecture/model choices and the trade-offs behind them
- ★★★★■ Coding fluency independent of the AI/ML narrative

No candidate resume supplied; this reflects only what interviewers historically probe in candidate experience, not this specific candidate's resume.

- ★★★★★ 01 Explain your current AI/ML project end to end and your exact contribution.
- ★★★★★ 02 Design or extend an ML/LLM system in production (ingestion through deployment and monitoring).
- ★★★★★ 03 Solve a coding/DSA problem and explain time/space complexity.
- ★★★★■ 04 Explain your RAG or retrieval design and the trade-offs you made.
- ★★★★■ 05 How would you debug or roll back a production ML failure?
- ★★★★■ 06 Why this model/architecture choice, and what trade-offs did you weigh?
- ★★■ 07 Tell me about a time you had more work than you could manage / got negative feedback (Googliness).

PREPARATION SIGNAL

recurrence, not chronology

- ↑ Coding / DSA (algorithm problem, complexity analysis)
- ↑ Production ML reasoning (deploy, debug, scale a live system)
- LLM / GenAI reasoning (RAG, fine-tuning, inference, agents)
- Current AI/ML project deep dive, own exact contribution
- RAG / retrieval / vector search

Signals reflect recurrence and interviewer emphasis across the candidate pool, not chronological change -- the pool is too small per company to claim a real time trend even though interview dates are known.

INTERVIEW-DERIVED ROLE FOCUS

- ★★★★★ Production ML reasoning
- ★★★★★ Coding / DSA fundamentals
- ★★★★■ LLM / GenAI, RAG, agents
- ★★★★■ System design at scale
- ★★■ Behavioral / cross-team collaboration

Exact JD not supplied; rankings reflect observed interview emphasis only.

REVISION PRIORITY

- ★★★★★ Project story: problem -> design -> your contribution -> impact
- ★★★★★ Coding/DSA warm-up (complexity, edge cases, clean code)
- ★★★★■ Production ML: debug/rollback/scale scenario
- ★★■ LLM/RAG design trade-offs
- ★★■ Googliness / behavioral answers ready

30-MINUTE REVISION ORDER

- 01 10 min Project story: problem -> design -> your contribution -> impact
- 02 7 min Coding/DSA warm-up (complexity, edge cases, clean code)
- 03 6 min Production ML: debug/rollback/scale scenario
- 04 5 min LLM/RAG design trade-offs
- 05 2 min Googliness / behavioral answers ready